

## SQL Queries

# The SELECT-FROM-WHERE Structure

```
SELECT <attributes>  
FROM <tables>  
WHERE <conditions>
```

From relational algebra:

- ▶ SELECT <attributes> corresponds to projection
- ▶ FROM <tables> specifies the table in parentheses in a relational algebra expression and joins
- ▶ WHERE <conditions> corresponds to selection

# Projection

$\pi_{\text{first\_name}, \text{last\_name}}(\text{author})$

```
mysql> select first_name, last_name from author;
```

first_name	last_name
John	McCarthy
Dennis	Ritchie
Ken	Thompson
Claude	Shannon
Alan	Turing
Alonzo	Church
Perry	White
Moshe	Vardi
Roy	Batty

```
9 rows in set (0.00 sec)
```

# Asterisk

```
mysql> select * from author;
+-----+-----+-----+
| author_id | first_name | last_name |
+-----+-----+-----+
|          1 | John      | McCarthy  |
|          2 | Dennis    | Ritchie   |
|          3 | Ken       | Thompson  |
|          4 | Claude    | Shannon   |
|          5 | Alan      | Turing    |
|          6 | Alonzo    | Church    |
|          7 | Perry     | White     |
|          8 | Moshe     | Vardi     |
|          9 | Roy       | Batty     |
+-----+-----+-----+
9 rows in set (0.00 sec)
```

Notice that with no condition on select, all rows returned.

# Select

$\sigma_{year=2012}(book)$

```
mysql> select * from book where year = 2012;
```

	book_id		book_title		month
					year
					editor
	7		AAAI		July
	8		NIPS		July

```
2 rows in set (0.00 sec)
```

# String Matching with LIKE

Our where condition can match a pattern with like. Use a % for wildcard, i.e., matching any character sequence.

Which publications have "Turing" in their titles?

```
select * from pub where title like 'Turing%';
```

pub_id	title	book_id
4	Turing Machines	4
5	Turing Test	5

2 rows in set (0.00 sec)

Note that strings are not case-sensitive.

# Joins

The `FROM` clause takes one or more source tables from the database and combines them into one (large) table using the `JOIN` operator. Three kinds of joins:

- ▶ `CROSS JOIN`
- ▶ `INNER JOIN`
- ▶ `OUTER JOIN`

Since DB designs are typically factored into many tables, the join is the most important part of a query.

# CROSS JOIN

A CROSS JOIN matches every row of the first table with every row of the second table. Think of a cross join as a cartesian product.

The general syntax for a cross join is:

```
SELECT <select_header> FROM <table1> CROSS JOIN <table2>
```

or

```
SELECT <select_header> FROM <table1>, <table2>
```



# CROSS JOIN EXAMPLE

```
mysql> select * from pub cross join book;
```

Pub_id	title	year	editor	book_id	book_id	book_title	month
1	LISP	1960	8	1	1	CACM	April
2	Unix	1960	8	2	1	CACM	April
3	Info Theory	1960	8	3	1	CACM	April
4	Turing Machines	1960	8	4	1	CACM	April
5	Turing Test	1960	8	5	1	CACM	April
6	Lambda Calculus	1960	8	6	1	CACM	April
1	LISP	1974	8	1	2	CACM	July
2	Unix	1974	8	2	2	CACM	July
3	Info Theory	1974	8	3	2	CACM	July
4	Turing Machines	1974	8	4	2	CACM	July
5	Turing Test	1974	8	5	2	CACM	July
6	Lambda Calculus	1974	8	6	2	CACM	July
1	LISP	1948	2	1	3	BST	July

# LIMITing Results

If we don't want many results to scroll past the bottom of the screen we can limit the number of results using a LIMIT clause.

```
mysql> select * from pub, book limit 3;
```

pub_id	title	book_id	book_id	book_title	month
year	editor				
1	LISP	1	1	CACM	April
1960	8				
2	Unix	2	1	CACM	April
1960	8				
3	Info Theory	3	1	CACM	April
1960	8				

3 rows in set (0.00 sec)

The general form of the LIMIT clause is LIMIT **start**, **count**, where **start** is the first row returned and **count** is the number of rows returned. If a single value is given, **start** assumes the value 0.

# Inner Joins

A simple inner join uses an ON condition.

```
mysql> select * from pub join book on pub.book_id = book.book_id;
```

pub_id	title	year	editor	book_id	book_id	book_title	month
1	LISP	1960	8	1	1	CACM	April
2	Unix	1974	8	2	2	CACM	July
3	Info Theory	1948	2	3	3	BST	July
4	Turing Machines	November 1936	7	4	4	LMS	
5	Turing Test	October 1950	NULL	5	5	Mind	
6	Lambda Calculus	1941	NULL	6	6	AMS	Month

6 rows in set (0.00 sec)

Notice that book\_id appears twice, because we get one from each source table. We can fix that ...

# Natural Joins

The USING clause, also called a natural join, equijoins on a like-named column from each table and includes the join column only once.

```
mysql> select * from pub join book using (book_id);
```

book_id	pub_id	title	book_title	month	year
1	8	LISP	CACM	April	1960
2	8	Unix	CACM	July	1974
3	2	Info Theory	BST	July	1948
4	7	Turing Machines	LMS	November	1936
5	NULL	Turing Test	Mind	October	1950
6	NULL	Lambda Calculus	AMS	Month	1941

6 rows in set (0.00 sec)

# Many to Many Relationships

A single author can write many publications, and a single publication can have many authors. This is a many-to-many relationship, which is modeled in relational databases with a relationship (or link or bridge) table.

```
CREATE TABLE IF NOT EXISTS author_pub (  
  author_id INTEGER NOT NULL REFERENCES author(author_id),  
  pub_id INTEGER NOT NULL REFERENCES publication(pub_id),  
  author_position INTEGER NOT NULL, -- first author, second, etc?  
  PRIMARY KEY (author_id, pub_id)  
);
```

author\_pub table links the author and pub tables

- ▶ author\_id and pub\_id are foreign keys to author and pub tables
- ▶ ~(author\_id, pub\_id) is composite key for the table

# Joining Multiple Tables

We can join all three tables by chaining join clauses:

```
mysql> select *  
-> from author join author_pub using (author_id)  
-> join pub using (pub_id);
```

pub_id	a_id	first_name	last_name	a_pos	title
1	1	John	McCarthy	1	LISP
2	2	Dennis	Ritchie	1	Unix
2	3	Ken	Thompson	2	Unix
3	4	Claude	Shannon	1	Info Theory
4	5	Alan	Turing	1	Turing
5	5	Alan	Turing	1	Turing Test
6	6	Alonzo	Church	1	Lambda

7 rows in set (0.01 sec)



# The SELECT Pipeline

The evaluation order of select clauses is approximately:

1. FROM <source\_tables> - Designates source tables and combining into one working table.
1. WHERE <filter\_expression> - Filters specific rows of working table
2. GROUP BY <grouping\_expressions> - Groups sets of rows in the working table based on column values
3. SELECT <select\_heading> - Defines the result set columns and (if applicable) grouping aggregates.
4. HAVING <filter\_expression> - Filters specific rows of the grouped table. Requires a GROUP BY
5. DISTINCT - Eliminates duplicate rows.
6. ~ORDER BY <ordering\_expressions> - Sorts the rows of the result set
7. OFFSET <count> - Skips over rows at the beginning of the result set. Requires a LIMIT.
8. LIMIT <count> - Limits the result set output to a specific number of rows.

Evaluation order determines what can be cross referenced in clauses.



# Aggregate Functions

Operate on groups of rows. Some common ones: COUNT, SUM, AVG

```
mysql> select count(*) from book;
+-----+
| count(*) |
+-----+
|          8 |
+-----+
```

There are 8 rows in the book table.

```
mysql> select count(editor) from book;
+-----+
| count(editor) |
+-----+
|              6 |
+-----+
```

Notice that COUNT doesn't count NULL values.

# GROUP BY

The GROUP BY clause groups rows in the working table by the values in the specified column(s) and collapses each group into a single row.

- ▶ We can apply an aggregate function to the resulting groups
- ▶ If we don't apply an aggregate function, only the last row of a group is returned.
  - ▶ Since rows within groups are in no particular order, failing to apply an aggregate function would essentially give us a random result.

# Aggregate Functions on Groups

Aggregate functions apply some function the to the rows grouped together by a GROUP BY clause.

How many papers did each author write?

```
mysql> select author_id, last_name, count(author_id)
-> from author join author_pub using (author_id)
-> join pub using (pub_id)
-> group by author_id;
```

author_id	last_name	count(author_id)
1	McCarthy	1
2	Ritchie	1
3	Thompson	1
4	Shannon	1
5	Turing	2
6	Church	1

Aggregate function is applied to column in GROUP BY.

# Simple Summation

Here are the data in the dorm table:

```
mysql> select * from dorm;
+-----+-----+-----+
| dorm_id | name      | spaces |
+-----+-----+-----+
|      1 | Armstrong |    124 |
|      2 | Brown     |    158 |
|      3 | Caldwell  |    158 |
+-----+-----+-----+
3 rows in set (0.00 sec)
```

What is the total capacity (number of spaces) for all dorms?

# SUM

To find the total capacity for all dorms, sum the spaces column:

```
mysql> select sum(spaces) from dorm;
+-----+
| sum(spaces) |
+-----+
|           440 |
+-----+
1 row in set (0.00 sec)
```

Or use a column alias in the select list to make output clearer:

```
mysql> select sum(spaces) as total_capacity from dorm;
+-----+
| total_capacity |
+-----+
|           440 |
+-----+
1 row in set (0.00 sec)
```

# Grouping and Counting

What is the occupancy of each dorm?

First, get a feel for the data:

```
mysql> select * from dorm join student using (dorm_id) order by  
      dorm.name;
```

dorm_id	name	spaces	student_id	name	gpa
1	Armstrong	124	1	Alice	3.60
1	Armstrong	124	2	Bob	2.70
1	Armstrong	124	3	Cheng	3.90
2	Brown	158	4	Dhruv	3.40
2	Brown	158	5	Ellie	4.00
2	Brown	158	6	Fong	2.30
3	Caldwell	158	7	Gerd	4.00
3	Caldwell	158	8	Hal	2.20
3	Caldwell	158	9	Isaac	2.00
3	Caldwell	158	10	Jacque	4.00

We can see that there are three groups of dorms in the result, which we could group by `dorm_id` or `dorm.name`.

# Dorm Occupancy

So we group by `dorm.name` and count the rows in each group.

```
mysql> select dorm.name as dorm_name, count(*) as occupancy  
       -> from dorm join student using (dorm_id)  
       -> group by dorm.name;
```

dorm_name	occupancy
Armstrong	3
Brown	3
Caldwell	4

3 rows in set (0.00 sec)

# Sorting, Aliasing, and Limiting

Who wrote the most publications?

```
mysql> select author_id, last_name, count(author_id) as pub_count
-> from author join author_pub using (author_id) join pub using
      (pub_id)
-> group by author_id
-> order by pub_count desc;
```

author_id	last_name	pub_count
5	Turing	2
1	McCarthy	1
2	Ritchie	1
6	Church	1
3	Thompson	1
4	Shannon	1

6 rows in set (0.00 sec)

Notice that we also used an alias so we could reference the count in the ORDER BY clause



# Limiting Results

If we want only the answer from the last query we can use LIMIT:  
Who wrote the most publications?

```
mysql> select author_id, last_name, count(author_id) as pub_count
-> from author join author_pub using (author_id) join pub using
      (pub_id)
-> group by author_id
-> order by pub_count desc
-> limit 1;
```

```
+-----+-----+-----+
| author_id | last_name | pub_count |
+-----+-----+-----+
|          5 | Turing    |          2 |
+-----+-----+-----+
1 row in set (0.00 sec)
```

# HAVING

In the previous query we got the top author by pub count. If we want all authors having a particular pub count, we can use a HAVING clause.

```
mysql> select author_id, last_name, count(author_id) as pub_count
-> from author join author_pub using (author_id)
-> join pub using (pub_id)
-> group by author_id
-> having pub_count = 1;
```

Author_id	last_name	pub_count
1	McCarthy	1
2	Ritchie	1
3	Thompson	1
4	Shannon	1
6	Church	1

We can use comparisons like  $<$ ,  $>$ . Notice that Turing is not in the result.

# HAVING vs. WHERE Conditions

Functionally HAVING and WHERE do the same thing: they filter-in tuples. The difference is where they are evaluated in the SELECT pipeline.

- ▶ WHERE is evaluated only after the FROM clause that selects the source tables, so WHERE clauses can only reference expressions that do not contain aggregate functions
- ▶ HAVING is evaluated after GROUP BY, and SELECT, so HAVING clauses can reference any result column

Be aware that rows filtered out by a WHERE clause will not be included in a GROUP BY clause.

# WHERE vs. HAVING Example

WHERE clause can't refer to column aliases and aggregates in the SELECT list or apply functions to groups created by GROUP BY clauses.

```
mysql> select author_id, last_name, count(author_id) as pub_count
-> from author natural join author_pub natural join pub
-> where pub_count = 1
-> group by author_id;
ERROR 1054 (42S22): Unknown column 'pub_count' in 'where clause'
```

HAVING can refer to select columns.

```
mysql> select author_id, last_name, count(author_id) as pub_count
-> from author natural join author_pub natural join pub
-> group by author_id
-> having pub_count = 1;
```

```
+-----+-----+-----+
| author_id | last_name | pub_count |
+-----+-----+-----+
|          1 | McCarthy  |          1 |
|          2 | Ritchie   |          1 |
|          3 | Thompson  |          1 |
|          4 | Shannon   |          1 |
|          6 | Church    |          1 |
+-----+-----+-----+
5 rows in set (0.00 sec)
```